

Machine Learning DevOps Engineer Nanodegree Program
Project: A Dynamic Risk Assessment System

Installation Instructions

Uncompress the file "starter-file.zip".

Install Conda: <https://docs.conda.io/projects/conda/en/latest/user-guide/install/index.html>

Create a Conda Enviroment with Python 3.7

```
conda create -n MLDOE python=3.7
```

```
conda activate MLDOE
```

Install the requirements:

```
pip install -r requirements.txt
```

File Structure

```
$ tree --dirsfirst
```

```
.
├── ingesteddata
│   ├── finaldata.csv
│   └── ingestedfiles.txt
├── models
│   ├── apireturns.txt
│   ├── confusionmatrix.png
│   ├── latestscore.txt
│   └── trainedmodel.pkl
├── practicedata
│   ├── dataset1.csv
│   └── dataset2.csv
├── practicemodels
│   ├── apireturns.txt
│   ├── confusionmatrix.png
│   ├── latestscore.txt
│   └── trainedmodel.pkl
├── production_deployment
│   ├── ingestedfiles.txt
│   ├── latestscore.txt
│   └── trainedmodel.pkl
├── sourcedata
│   ├── dataset3.csv
│   └── dataset4.csv
├── testdata
│   └── testdata.csv
├── apicalls.py
├── app.py
├── clear_results.py
├── config.json
├── config.json.backup1
├── config.json.backup2
├── crontab.file
├── deployment.py
├── diagnostics.py
├── execute_full_process_twice.sh
├── fullprocess.py
├── ingestion.py
├── reporting.py
├── REPORT.odt
├── REPORT.pdf
├── requirements.txt
└── scoring.py
```

```
├─ training.py
├─ use_configuration_1.sh
├─ use_configuration_2.sh
└─ wsgi.py
```

7 directories, 39 files

Example to Execute (with drift_must_improve_score = False)

Given that the practice model gives a better F1-score than the final model, it was necessary to tweak the Python script fullprocess.py in order to execute the final process twice and to obtain 2 confusion matrices and 2 apireturns.txt files:

```
drift_must_improve_score = False

def check_for_model_drift():
    score0 = scoring.read_f1_score(load_config()['prod_deployment_path'],
    'latestscore.txt')
    print(f'score0={score0}')
    os.system('python training.py')
    os.system('python scoring.py')
    score1 = scoring.read_f1_score(load_config()['output_model_path'], 'latestscore.txt')
    print(f'score0={score0}, score1={score1}')
    return (score1 > score0) if drift_must_improve_score else abs(score1 - score0) >
0.001
```

In this first execution of execute_full_process_twice.sh, drift can be positive and negative, but not zero: (drift_must_improve_score = False) In the second execution of execute_full_process_twice.sh, drift can be only positive, greater than zero: (drift_must_improve_score = True)

Important Note: The Server Side (python app.py) must be executed before executing the script "execute_full_process_twice.sh". Otherwise the RESTful API won't be available.

\$./execute_full_process_twice.sh

```
===== EXECUTING execute_full_process_twice.sh =====
echo "===== EXECUTING execute_full_process_twice.sh ====="
cat execute_full_process_twice.sh
echo "===== EXECUTING python clear_results.py ====="
python clear_results.py
echo "===== EXECUTING python fullprocess.py (FOR THE FIRST TIME) ====="
python fullprocess.py
echo "===== EXECUTING ./use_configuration_2.sh ====="
./use_configuration_2.sh
echo "===== EXECUTING python fullprocess.py (FOR THE SECOND TIME) ====="
python fullprocess.py
===== EXECUTING python clear_results.py =====
===== EXECUTING python fullprocess.py (FOR THE FIRST TIME) =====
new_csv_files=['practicedata/dataset1.csv', 'practicedata/dataset2.csv']
-----
There are new CSV files. Running the script ingestion.py.
-----
===== RUNNING PYTHON SCRIPT ingestion.py =====
csv_files=['practicedata/dataset1.csv', 'practicedata/dataset2.csv']
practicedata/dataset1.csv
  corporation  lastmonth_activity  lastyear_activity  number_of_employees  exited
0          nciw             100             1359                1          0
1          lsid              68              282               14          0
2          pwls              71              949               40          1
3          bqlx             686             3782              103          0
4          zmei              45              655                7          0
5          wosl               0               18               21          1
6          xcvb             189              961               18          1
7          dfgh              16             1028               33          0
```

8	ngrd	9	45	1	1
9	xful	0	67	14	1
10	kshe	48	986	22	1
11	qqqq	52	650	11	1
12	corp	1090	2452	9	0
13	ekci	6	88	90	1
14	dosk	99	390	99	1
15	endi	75	800	81	1
16	gudj	255	1687	2	0

practicedata/dataset2.csv

	corporation	lastmonth_activity	lastyear_activity	number_of_employees	exited
0	abcd	78	1024	12	1
1	asdf	14	2145	20	0
2	xyzz	182	3891	35	0
3	acme	101	10983	2	1
4	qwer	0	118	42	1
5	tyui	929	1992	1	0
6	zxcv	19	455	8	1
7	hjkl	94	868	3	1
8	lmno	81	1401	10	0
9	qqqq	52	650	11	1
10	corp	1090	2452	9	0
11	ekci	6	88	90	1
12	dosk	99	390	99	1
13	endi	75	800	81	1
14	gudj	255	1687	2	0
15	wosl	0	18	21	1
16	xcvb	189	961	18	1
17	dfgh	16	1028	33	0
18	ngrd	9	45	1	1

ingesteddata/finaldata.csv

	corporation	lastmonth_activity	lastyear_activity	number_of_employees	exited
0	nciw	100	1359	1	0
1	lsid	68	282	14	0
2	pwls	71	949	40	1
3	bqlx	686	3782	103	0
4	zmei	45	655	7	0
5	wosl	0	18	21	1
6	xcvb	189	961	18	1
7	dfgh	16	1028	33	0
8	ngrd	9	45	1	1
9	xful	0	67	14	1
10	kshe	48	986	22	1
11	qqqq	52	650	11	1
12	corp	1090	2452	9	0
13	ekci	6	88	90	1
14	dosk	99	390	99	1
15	endi	75	800	81	1
16	gudj	255	1687	2	0
17	abcd	78	1024	12	1
18	asdf	14	2145	20	0
19	xyzz	182	3891	35	0
20	acme	101	10983	2	1
21	qwer	0	118	42	1
22	tyui	929	1992	1	0
23	zxcv	19	455	8	1
24	hjkl	94	868	3	1
25	lmno	81	1401	10	0

Record of ingestion was saved in file "ingesteddata/ingestedfiles.txt".

===== TESTING FOR DRIFT =====

score0=0.0

csv_file=ingesteddata/finaldata.csv

```
X=[[ 100 1359    1]
 [  68  282   14]
 [  71  949   40]
 [ 686 3782  103]
 [  45  655    7]
```

```

[ 0 18 21]
[ 189 961 18]
[ 16 1028 33]
[ 9 45 1]
[ 0 67 14]
[ 48 986 22]
[ 52 650 11]
[ 1090 2452 9]
[ 6 88 90]
[ 99 390 99]
[ 75 800 81]
[ 255 1687 2]
[ 78 1024 12]
[ 14 2145 20]
[ 182 3891 35]
[ 101 10983 2]
[ 0 118 42]
[ 929 1992 1]
[ 19 455 8]
[ 94 868 3]
[ 81 1401 10]]
Y=[0 0 1 0 0 1 1 0 1 1 1 0 1 1 1 0 1 0 0 1 1 0 1 1 0]
Model was saved in the file "practicemodels/trainedmodel.pkl".
Test data loaded from file "testdata/testdata.csv".
X=[[ 234 3 10]
[ 14 2145 99]
[ 34 333 1000]
[ 101 12346 2]
[ 0 675 25]]
Y=[1 0 0 1 1]
Model was loaded from file "practicemodels/trainedmodel.pkl".
f1 score: 0.5714285714285715
F1-score 0.5714285714285715 saved in file "practicemodels/latestscore.txt".
score0=0.0, score1=0.5714285714285715

-----
There is drift. Running the scripts deployment.py, apicalls.py, and reporting.py.
-----

===== RUNNING PYTHON SCRIPT deployment.py =====
File "practicemodels/trainedmodel.pkl" was successfully copied to
"production_deployment".
File "practicemodels/latestscore.txt" was successfully copied to "production_deployment".
File "ingesteddata/ingestedfiles.txt" was successfully copied to "production_deployment".
===== RUNNING PYTHON SCRIPT apicalls.py =====
INPUT 1: {'location': 'testdata/testdata.csv'}
OUTPUT 1: {"predictions": [0, 1, 1, 1, 1]}
OUTPUT 2: {'f1_score': 0.5714285714285715}
OUTPUT 3: {'exited': {'mean': 0.5769230769230769, 'median': 1.0, 'stdev':
0.4940474068717357}, 'lastmonth_activity': {'mean': 165.65384615384616, 'median': 73.0,
'stdev': 278.5174959713127}, 'lastyear_activity': {'mean': 1502.923076923077, 'median':
955.0, 'stdev': 2150.065274913888}, 'number_of_employees': {'mean': 26.884615384615383,
'median': 14.0, 'stdev': 30.745014509018585}}
OUTPUT 4: {'na_percentages': [0.0, 0.0, 0.0, 0.0, 0.0], 'outdated_packages': {'Pillow':
['8.1.0', '8.1.0', '8.1.2'], 'pandas': ['1.2.2', '1.2.2', '1.2.3']}, 'times':
[1.0436301231384277, 0.46016383171081543]}
The API returns were saved in file "practicemodels/apireturns.txt".
===== RUNNING PYTHON SCRIPT reporting.py =====
Test data loaded from file "testdata/testdata.csv".
X=[[ 234 3 10]
[ 14 2145 99]
[ 34 333 1000]
[ 101 12346 2]
[ 0 675 25]]
Y=[1 0 0 1 1]
Predictions:
[0, 1, 1, 1, 1]
tn=0, fp=2, fn=1, tp=2
Confusion matrix plot saved to the file "practicemodels/confusionmatrix.png".

```

```
===== EXECUTING ./use_configuration_2.sh =====
===== EXECUTING python fullprocess.py (FOR THE SECOND TIME) =====
new_csv_files=['sourcedata/dataset3.csv', 'sourcedata/dataset4.csv']
```

There are new CSV files. Running the script ingestion.py.

```
===== RUNNING PYTHON SCRIPT ingestion.py =====
```

```
csv_files=['sourcedata/dataset3.csv', 'sourcedata/dataset4.csv']
```

```
sourcedata/dataset3.csv
```

	corporation	lastmonth_activity	lastyear_activity	number_of_employees	exited
0	nciw	45	0	99	1
1	lsid	36	234	541	0
2	pwls	23	555	23	0
3	bqlx	15	11	190	1
4	zmei	100	2929	999	1
5	wosl	2	1	1359	0
6	xcvb	0	14	282	1
7	dfgh	500	40	949	0
8	ngrd	1234	103	3782	1
9	xful	98765	7	655	0
10	kshe	34	2345	18	1

```
sourcedata/dataset4.csv
```

	corporation	lastmonth_activity	lastyear_activity	number_of_employees	exited
0	abcd	99	871	3	0
1	asdf	1243	0	10	0
2	xyzz	0	25	11	1
3	acme	813	129	9	1
4	qwer	2989	9982	90	0
5	tyui	395	190	99	0
6	zxcv	19028	999	81	0
7	hjkl	345	78	2	1
8	lmno	1024	14	3	1
9	qqqq	2145	182	110	1
10	corp	3891	101	998	0
11	ekci	10983	0	1200	0
12	dosk	118	929	81	1
13	endi	1992	19	2	0
14	gudj	455	94	298	1

```
ingesteddata/finaldata.csv
```

	corporation	lastmonth_activity	lastyear_activity	number_of_employees	exited
0	nciw	45	0	99	1
1	lsid	36	234	541	0
2	pwls	23	555	23	0
3	bqlx	15	11	190	1
4	zmei	100	2929	999	1
5	wosl	2	1	1359	0
6	xcvb	0	14	282	1
7	dfgh	500	40	949	0
8	ngrd	1234	103	3782	1
9	xful	98765	7	655	0
10	kshe	34	2345	18	1
11	abcd	99	871	3	0
12	asdf	1243	0	10	0
13	xyzz	0	25	11	1
14	acme	813	129	9	1
15	qwer	2989	9982	90	0
16	tyui	395	190	99	0
17	zxcv	19028	999	81	0
18	hjkl	345	78	2	1
19	lmno	1024	14	3	1
20	qqqq	2145	182	110	1
21	corp	3891	101	998	0
22	ekci	10983	0	1200	0
23	dosk	118	929	81	1
24	endi	1992	19	2	0
25	gudj	455	94	298	1

Record of ingestion was saved in file "ingesteddata/ingestedfiles.txt".

===== TESTING FOR DRIFT =====

score0=0.5714285714285715

csv_file=ingesteddata/finaldata.csv

X=[[45 0 99]

[36 234 541]

[23 555 23]

[15 11 190]

[100 2929 999]

[2 1 1359]

[0 14 282]

[500 40 949]

[1234 103 3782]

[98765 7 655]

[34 2345 18]

[99 871 3]

[1243 0 10]

[0 25 11]

[813 129 9]

[2989 9982 90]

[395 190 99]

[19028 999 81]

[345 78 2]

[1024 14 3]

[2145 182 110]

[3891 101 998]

[10983 0 1200]

[118 929 81]

[1992 19 2]

[455 94 298]]

Y=[1 0 0 1 1 0 1 0 1 0 1 0 0 1 1 0 0 0 1 1 1 0 0 1 0 1]

Model was saved in the file "models/trainedmodel.pkl".

Test data loaded from file "testdata/testdata.csv".

X=[[234 3 10]

[14 2145 99]

[34 333 1000]

[101 12346 2]

[0 675 25]]

Y=[1 0 0 1 1]

Model was loaded from file "models/trainedmodel.pkl".

f1 score: 0.3333333333333333

F1-score 0.3333333333333333 saved in file "models/latestscore.txt".

score0=0.5714285714285715, score1=0.3333333333333333

There is drift. Running the scripts deployment.py, apicalls.py, and reporting.py.

===== RUNNING PYTHON SCRIPT deployment.py =====

File "models/trainedmodel.pkl" was successfully copied to "production_deployment".

File "models/latestscore.txt" was successfully copied to "production_deployment".

File "ingesteddata/ingestedfiles.txt" was successfully copied to "production_deployment".

===== RUNNING PYTHON SCRIPT apicalls.py =====

INPUT 1: {'location': 'testdata/testdata.csv'}

OUTPUT 1: {"predictions": [0, 1, 1, 0, 1]}

OUTPUT 2: {'f1_score': 0.3333333333333333}

OUTPUT 3: {'exited': {'mean': 0.5, 'median': 0.5, 'stdev': 0.5}, 'lastmonth_activity': {'mean': 5625.923076923077, 'median': 425.0, 'stdev': 19067.170236829497}, 'lastyear_activity': {'mean': 763.5384615384615, 'median': 97.5, 'stdev': 1977.4481658240022}, 'number_of_employees': {'mean': 457.46153846153845, 'median': 99.0, 'stdev': 785.0576280057029}}

OUTPUT 4: {'na_percentages': [0.0, 0.0, 0.0, 0.0, 0.0], 'outdated_packages': {'Pillow': ['8.1.0', '8.1.0', '8.1.2'], 'pandas': ['1.2.2', '1.2.2', '1.2.3']}, 'times': [1.0206544399261475, 0.46021509170532227]}

The API returns were saved in file "models/apireturns.txt".

===== RUNNING PYTHON SCRIPT reporting.py =====

Test data loaded from file "testdata/testdata.csv".

X=[[234 3 10]

[14 2145 99]

[34 333 1000]

```

[ 101 12346    2]
[    0    675   25]]
Y=[1 0 0 1 1]
Predictions:
[0, 1, 1, 0, 1]
tn=0, fp=2, fn=2, tp=1
Confusion matrix plot saved to the file "models/confusionmatrix.png".

```

Server Side

\$ python app.py

```

* Serving Flask app "app" (lazy loading)
* Environment: production
  WARNING: This is a development server. Do not use it in a production deployment.
  Use a production WSGI server instead.
* Debug mode: on
* Running on http://0.0.0.0:8000/ (Press CTRL+C to quit)
* Restarting with stat
* Debugger is active!
* Debugger PIN: 166-233-122
X=[[ 234    3   10]
 [ 14 2145   99]
 [ 34 333 1000]
 [ 101 12346    2]
 [    0    675   25]]
Y=[1 0 0 1 1]
127.0.0.1 - - [16/Mar/2021 05:27:00] "POST /prediction HTTP/1.1" 200 -
Test data loaded from file "testdata/testdata.csv".
X=[[ 234    3   10]
 [ 14 2145   99]
 [ 34 333 1000]
 [ 101 12346    2]
 [    0    675   25]]
Y=[1 0 0 1 1]
Model was loaded from file "practicemodels/trainedmodel.pkl".
f1 score: 0.5714285714285715
F1-score 0.5714285714285715 saved in file "practicemodels/latestscore.txt".
127.0.0.1 - - [16/Mar/2021 05:27:00] "GET /scoring HTTP/1.1" 200 -
Test data loaded from file "ingesteddata/finaldata.csv".
127.0.0.1 - - [16/Mar/2021 05:27:00] "GET /summarystats HTTP/1.1" 200 -
Test data loaded from file "ingesteddata/finaldata.csv".
WARNING: pip is being invoked by an old script wrapper. This will fail in a future
version of pip.
Please see https://github.com/pypa/pip/issues/5599 for advice on fixing the underlying
issue.
To avoid this problem you can invoke Python with '-m pip' instead of running pip
directly.
WARNING: pip is being invoked by an old script wrapper. This will fail in a future
version of pip.
Please see https://github.com/pypa/pip/issues/5599 for advice on fixing the underlying
issue.
To avoid this problem you can invoke Python with '-m pip' instead of running pip
directly.
127.0.0.1 - - [16/Mar/2021 05:27:11] "GET /diagnostics HTTP/1.1" 200 -
X=[[ 234    3   10]
 [ 14 2145   99]
 [ 34 333 1000]
 [ 101 12346    2]
 [    0    675   25]]
Y=[1 0 0 1 1]
127.0.0.1 - - [16/Mar/2021 05:27:19] "POST /prediction HTTP/1.1" 200 -
Test data loaded from file "testdata/testdata.csv".
X=[[ 234    3   10]
 [ 14 2145   99]
 [ 34 333 1000]

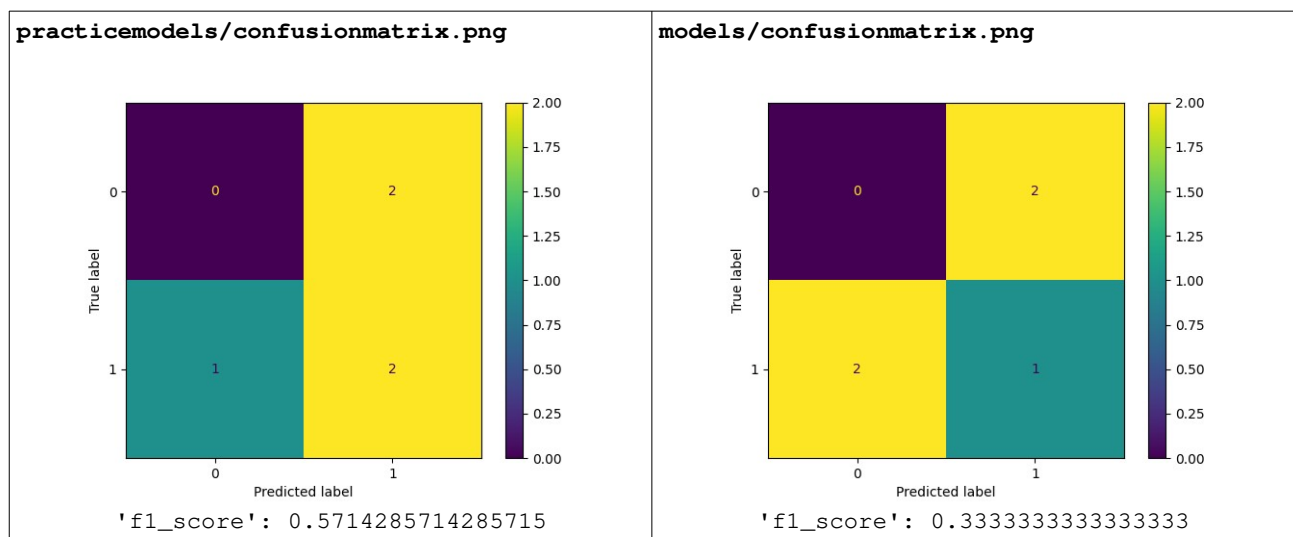
```

```

[ 101 12346    2]
[   0    675   25]]
Y=[1 0 0 1 1]
Model was loaded from file "models/trainedmodel.pkl".
f1 score: 0.3333333333333333
F1-score 0.3333333333333333 saved in file "models/latestscore.txt".
127.0.0.1 - - [16/Mar/2021 05:27:19] "GET /scoring HTTP/1.1" 200 -
Test data loaded from file "ingesteddata/finaldata.csv".
127.0.0.1 - - [16/Mar/2021 05:27:19] "GET /summarystats HTTP/1.1" 200 -
Test data loaded from file "ingesteddata/finaldata.csv".
WARNING: pip is being invoked by an old script wrapper. This will fail in a future
version of pip.
Please see https://github.com/pypa/pip/issues/5599 for advice on fixing the underlying
issue.
To avoid this problem you can invoke Python with '-m pip' instead of running pip
directly.
WARNING: pip is being invoked by an old script wrapper. This will fail in a future
version of pip.
Please see https://github.com/pypa/pip/issues/5599 for advice on fixing the underlying
issue.
To avoid this problem you can invoke Python with '-m pip' instead of running pip
directly.
127.0.0.1 - - [16/Mar/2021 05:27:30] "GET /diagnostics HTTP/1.1" 200 -

```

API Returns and Confusion Matrices



\$ cat practicemodels/apireturns.txt

```

INPUT 1: {'location': 'testdata/testdata.csv'}
OUTPUT 1: {"predictions": [0, 1, 1, 1, 1]}
OUTPUT 2: {'f1_score': 0.5714285714285715}
OUTPUT 3: {'exited': {'mean': 0.5769230769230769, 'median': 1.0, 'stdev':
0.4940474068717357}, 'lastmonth_activity': {'mean': 165.65384615384616, 'median': 73.0,
'stdev': 278.5174959713127}, 'lastyear_activity': {'mean': 1502.923076923077, 'median':
955.0, 'stdev': 2150.065274913888}, 'number_of_employees': {'mean': 26.884615384615383,
'median': 14.0, 'stdev': 30.745014509018585}}
OUTPUT 4: {'na_percentages': [0.0, 0.0, 0.0, 0.0, 0.0], 'outdated_packages': {'Pillow':
['8.1.0', '8.1.0', '8.1.2'], 'pandas': ['1.2.2', '1.2.2', '1.2.3']}, 'times':
[1.0436301231384277, 0.46016383171081543]}

```

\$ cat models/apireturns.txt

```

INPUT 1: {'location': 'testdata/testdata.csv'}
OUTPUT 1: {"predictions": [0, 1, 1, 0, 1]}
OUTPUT 2: {'f1_score': 0.3333333333333333}
OUTPUT 3: {'exited': {'mean': 0.5, 'median': 0.5, 'stdev': 0.5}, 'lastmonth_activity':
{'mean': 5625.923076923077, 'median': 425.0, 'stdev': 19067.170236829497},

```



```
'lastyear_activity': {'mean': 763.5384615384615, 'median': 97.5, 'stdev':
1977.4481658240022}, 'number_of_employees': {'mean': 457.46153846153845, 'median': 99.0,
'stdev': 785.0576280057029}}
OUTPUT 4: {'na_percentages': [0.0, 0.0, 0.0, 0.0, 0.0], 'outdated_packages': {'Pillow':
['8.1.0', '8.1.0', '8.1.2'], 'pandas': ['1.2.2', '1.2.2', '1.2.3']}, 'times':
[1.0206544399261475, 0.46021509170532227]}
```

Example to Execute (with drift_must_improve_score = True)

In the first execution of execute_full_process_twice.sh, drift could be positive and negative, but not zero: (drift_must_improve_score = False) In this second execution of execute_full_process_twice.sh, drift can be only positive, greater than zero: (drift_must_improve_score = True). As a result, only in the first execution of "python fullprocess.py", drift is generated. And drift is not generated in the second execution of "python fullprocess.py". In other words, there is only 1 apireturns.txt file and 1 confusion matrix, not 2 pairs.

```
drift_must_improve_score = True

def check_for_model_drift():
    score0 = scoring.read_fl_score(load_config()['prod_deployment_path'],
'latestscore.txt')
    print(f'score0={score0}')
    os.system('python training.py')
    os.system('python scoring.py')
    score1 = scoring.read_fl_score(load_config()['output_model_path'], 'latestscore.txt')
    print(f'score0={score0}, score1={score1}')
    return (score1 > score0) if drift_must_improve_score else abs(score1 - score0) >
0.001
```

```
$ ./execute_full_process_twice.sh
===== EXECUTING execute_full_process_twice.sh =====
echo "===== EXECUTING execute_full_process_twice.sh ====="
cat execute_full_process_twice.sh
echo "===== EXECUTING python clear_results.py ====="
python clear_results.py
echo "===== EXECUTING python fullprocess.py (FOR THE FIRST TIME) ====="
python fullprocess.py
echo "===== EXECUTING ./use_configuration_2.sh ====="
./use_configuration_2.sh
echo "===== EXECUTING python fullprocess.py (FOR THE SECOND TIME) ====="
python fullprocess.py
===== EXECUTING python clear_results.py =====
===== EXECUTING python fullprocess.py (FOR THE FIRST TIME) =====
new_csv_files=['practicedata/dataset1.csv', 'practicedata/dataset2.csv']
-----
There are new CSV files. Running the script ingestion.py.
-----
===== RUNNING PYTHON SCRIPT ingestion.py =====
csv_files=['practicedata/dataset1.csv', 'practicedata/dataset2.csv']
practicedata/dataset1.csv
  corporation  lastmonth_activity  lastyear_activity  number_of_employees  exited
0          nciw             100             1359              1          0
1          lsid              68              282             14          0
2          pwls              71              949             40          1
3          bqlx             686             3782            103          0
4          zmei              45              655              7          0
5          wosl              0               18             21          1
6          xcvb             189              961             18          1
7          dfgh             16             1028             33          0
8          ngrd              9               45              1          1
9          xful              0               67             14          1
10         kshe             48              986             22          1
11         qqqq             52              650             11          1
```

12	corp	1090	2452	9	0
13	ekci	6	88	90	1
14	dosk	99	390	99	1
15	endi	75	800	81	1
16	gudj	255	1687	2	0

practicedata/dataset2.csv

	corporation	lastmonth_activity	lastyear_activity	number_of_employees	exited
0	abcd	78	1024	12	1
1	asdf	14	2145	20	0
2	xyzz	182	3891	35	0
3	acme	101	10983	2	1
4	qwer	0	118	42	1
5	tyui	929	1992	1	0
6	zxcv	19	455	8	1
7	hjkl	94	868	3	1
8	lmno	81	1401	10	0
9	qqqq	52	650	11	1
10	corp	1090	2452	9	0
11	ekci	6	88	90	1
12	dosk	99	390	99	1
13	endi	75	800	81	1
14	gudj	255	1687	2	0
15	wosl	0	18	21	1
16	xcvb	189	961	18	1
17	dfgh	16	1028	33	0
18	ngrd	9	45	1	1

ingesteddata/finaldata.csv

	corporation	lastmonth_activity	lastyear_activity	number_of_employees	exited
0	nciw	100	1359	1	0
1	lsid	68	282	14	0
2	pwls	71	949	40	1
3	bqlx	686	3782	103	0
4	zmei	45	655	7	0
5	wosl	0	18	21	1
6	xcvb	189	961	18	1
7	dfgh	16	1028	33	0
8	ngrd	9	45	1	1
9	xful	0	67	14	1
10	kshe	48	986	22	1
11	qqqq	52	650	11	1
12	corp	1090	2452	9	0
13	ekci	6	88	90	1
14	dosk	99	390	99	1
15	endi	75	800	81	1
16	gudj	255	1687	2	0
17	abcd	78	1024	12	1
18	asdf	14	2145	20	0
19	xyzz	182	3891	35	0
20	acme	101	10983	2	1
21	qwer	0	118	42	1
22	tyui	929	1992	1	0
23	zxcv	19	455	8	1
24	hjkl	94	868	3	1
25	lmno	81	1401	10	0

Record of ingestion was saved in file "ingesteddata/ingestedfiles.txt".

===== TESTING FOR DRIFT =====

score0=0.0

csv_file=ingesteddata/finaldata.csv

```
X=[[ 100 1359    1]
 [  68  282   14]
 [  71  949   40]
 [ 686 3782  103]
 [  45  655    7]
 [   0   18   21]
 [ 189  961   18]
 [  16 1028   33]
 [   9   45    1]
```

```

[ 0 67 14]
[ 48 986 22]
[ 52 650 11]
[ 1090 2452 9]
[ 6 88 90]
[ 99 390 99]
[ 75 800 81]
[ 255 1687 2]
[ 78 1024 12]
[ 14 2145 20]
[ 182 3891 35]
[ 101 10983 2]
[ 0 118 42]
[ 929 1992 1]
[ 19 455 8]
[ 94 868 3]
[ 81 1401 10]]
Y=[0 0 1 0 0 1 1 0 1 1 1 0 1 1 0 1 0 0 1 1 0 1 1 0]
Model was saved in the file "practicemodels/trainedmodel.pkl".
Test data loaded from file "testdata/testdata.csv".
X=[[ 234 3 10]
[ 14 2145 99]
[ 34 333 1000]
[ 101 12346 2]
[ 0 675 25]]
Y=[1 0 0 1 1]
Model was loaded from file "practicemodels/trainedmodel.pkl".
f1 score: 0.5714285714285715
F1-score 0.5714285714285715 saved in file "practicemodels/latestscore.txt".
score0=0.0, score1=0.5714285714285715

-----
There is drift. Running the scripts deployment.py, apicalls.py, and reporting.py.
-----

===== RUNNING PYTHON SCRIPT deployment.py =====
File "practicemodels/trainedmodel.pkl" was successfully copied to
"production_deployment".
File "practicemodels/latestscore.txt" was successfully copied to "production_deployment".
File "ingesteddata/ingestedfiles.txt" was successfully copied to "production_deployment".
===== RUNNING PYTHON SCRIPT apicalls.py =====
INPUT 1: {'location': 'testdata/testdata.csv'}
OUTPUT 1: {"predictions": [0, 1, 1, 1, 1]}
OUTPUT 2: {'f1_score': 0.5714285714285715}
OUTPUT 3: {'exited': {'mean': 0.5769230769230769, 'median': 1.0, 'stdev':
0.4940474068717357}, 'lastmonth_activity': {'mean': 165.65384615384616, 'median': 73.0,
'stdev': 278.5174959713127}, 'lastyear_activity': {'mean': 1502.923076923077, 'median':
955.0, 'stdev': 2150.065274913888}, 'number_of_employees': {'mean': 26.884615384615383,
'median': 14.0, 'stdev': 30.745014509018585}}
OUTPUT 4: {'na_percentages': [0.0, 0.0, 0.0, 0.0, 0.0], 'outdated_packages': {'Pillow':
['8.1.0', '8.1.0', '8.1.2'], 'pandas': ['1.2.2', '1.2.2', '1.2.3']}, 'times':
[1.035158634185791, 0.4606971740722656]}
The API returns were saved in file "practicemodels/apireturns.txt".
===== RUNNING PYTHON SCRIPT reporting.py =====
Test data loaded from file "testdata/testdata.csv".
X=[[ 234 3 10]
[ 14 2145 99]
[ 34 333 1000]
[ 101 12346 2]
[ 0 675 25]]
Y=[1 0 0 1 1]
Predictions:
[0, 1, 1, 1, 1]
tn=0, fp=2, fn=1, tp=2
Confusion matrix plot saved to the file "practicemodels/confusionmatrix.png".
===== EXECUTING ./use_configuration_2.sh =====
===== EXECUTING python fullprocess.py (FOR THE SECOND TIME) =====
new_csv_files=['sourcedata/dataset3.csv', 'sourcedata/dataset4.csv']
-----

```

There are new CSV files. Running the script ingestion.py.

===== RUNNING PYTHON SCRIPT ingestion.py =====

csv_files=['sourcedata/dataset3.csv', 'sourcedata/dataset4.csv']

sourcedata/dataset3.csv

	corporation	lastmonth_activity	lastyear_activity	number_of_employees	exited
0	nciw	45	0	99	1
1	lsid	36	234	541	0
2	pwls	23	555	23	0
3	bqlx	15	11	190	1
4	zmei	100	2929	999	1
5	wosl	2	1	1359	0
6	xcvb	0	14	282	1
7	dfgh	500	40	949	0
8	ngrd	1234	103	3782	1
9	xful	98765	7	655	0
10	kshe	34	2345	18	1

sourcedata/dataset4.csv

	corporation	lastmonth_activity	lastyear_activity	number_of_employees	exited
0	abcd	99	871	3	0
1	asdf	1243	0	10	0
2	xyzz	0	25	11	1
3	acme	813	129	9	1
4	qwer	2989	9982	90	0
5	tyui	395	190	99	0
6	zxcv	19028	999	81	0
7	hjkl	345	78	2	1
8	lmno	1024	14	3	1
9	qqqq	2145	182	110	1
10	corp	3891	101	998	0
11	ekci	10983	0	1200	0
12	dosk	118	929	81	1
13	endi	1992	19	2	0
14	gudj	455	94	298	1

ingesteddata/finaldata.csv

	corporation	lastmonth_activity	lastyear_activity	number_of_employees	exited
0	nciw	45	0	99	1
1	lsid	36	234	541	0
2	pwls	23	555	23	0
3	bqlx	15	11	190	1
4	zmei	100	2929	999	1
5	wosl	2	1	1359	0
6	xcvb	0	14	282	1
7	dfgh	500	40	949	0
8	ngrd	1234	103	3782	1
9	xful	98765	7	655	0
10	kshe	34	2345	18	1
11	abcd	99	871	3	0
12	asdf	1243	0	10	0
13	xyzz	0	25	11	1
14	acme	813	129	9	1
15	qwer	2989	9982	90	0
16	tyui	395	190	99	0
17	zxcv	19028	999	81	0
18	hjkl	345	78	2	1
19	lmno	1024	14	3	1
20	qqqq	2145	182	110	1
21	corp	3891	101	998	0
22	ekci	10983	0	1200	0
23	dosk	118	929	81	1
24	endi	1992	19	2	0
25	gudj	455	94	298	1

Record of ingestion was saved in file "ingesteddata/ingestedfiles.txt".

===== TESTING FOR DRIFT =====

score0=0.5714285714285715

csv_file=ingesteddata/finaldata.csv

X=[[45 0 99]

```
[ 36 234 541]
[ 23 555 23]
[ 15 11 190]
[ 100 2929 999]
[ 2 1 1359]
[ 0 14 282]
[ 500 40 949]
[ 1234 103 3782]
[98765 7 655]
[ 34 2345 18]
[ 99 871 3]
[ 1243 0 10]
[ 0 25 11]
[ 813 129 9]
[ 2989 9982 90]
[ 395 190 99]
[19028 999 81]
[ 345 78 2]
[ 1024 14 3]
[ 2145 182 110]
[ 3891 101 998]
[10983 0 1200]
[ 118 929 81]
[ 1992 19 2]
[ 455 94 298]]
Y=[1 0 0 1 1 0 1 0 1 0 1 0 0 1 1 0 0 0 1 1 1 0 0 1 0 1]
Model was saved in the file "models/trainedmodel.pkl".
Test data loaded from file "testdata/testdata.csv".
X=[[ 234 3 10]
[ 14 2145 99]
[ 34 333 1000]
[ 101 12346 2]
[ 0 675 25]]
Y=[1 0 0 1 1]
Model was loaded from file "models/trainedmodel.pkl".
f1 score: 0.3333333333333333
F1-score 0.3333333333333333 saved in file "models/latestscore.txt".
score0=0.5714285714285715, score1=0.3333333333333333
There is no drift. Process ended.
```
